The Smithlab DNA Methylation Data Analysis Pipeline (MethPipe)

Qiang Song   Benjamin Decato   Michael Kessler   Fang Fang   Jenny Qu
Tyler Garvin   Meng Zhou   Andrew Smith

July 24, 2014

The `methpipe` software package is a comprehensive pipeline and set of tools for analyzing whole genome bisulfite sequencing data (BS-seq). This manual explains the stages in our pipeline, how to use the analysis tools, and how to modify the pipeline for your specific context.

Contents

1 Assumptions

Our pipeline was designed to run in a cluster computing context, with many processing nodes available, and a job submission system like PBS or SGE. Much of this analysis is computationally intensive. We assume that individual nodes will have several GB of memory available for processing. Typically the data we deal with amounts to a minimum of 100GB for a mammalian methylome at 10x coverage. Intermediate files may cause this amount to more than double during execution of the pipeline, and likely at the end of the pipeline the total size of files will amount to almost double the size of the raw data.

Users are assumed to be quite familiar with UNIX/Linux and related concepts (*e.g.* building software from source, using the command line, shell environment variables, etc.).

It is also critical that users are familiar with BS-seq experiments, especially the bisulfite conversion reaction, and how this affects what we observe in the sequenced reads. This is especially important if paired-end sequencing is used. If you do not understand these concepts, you will likely run into major problems trying to customize our pipeline.

2 Methylome construction

2.1 Mapping reads

During bisulfite treatment, unmethylated cytosines in the original DNA sequences are converted to uracils, which are then incorporated as thymines (T) during PCR amplification. These PCR products are referred to as T-rich sequences as a result of their high thymine constitution. With paired-end sequencing experiments, the compliments of these T-rich sequences are also sequenced. These complimentary sequences have high adenosine (A) constitution (A is the complimentary base pair of T), and are referred to as A-rich sequences. Mapping consists of finding sequence similarity, based on context specific criteria, between these short sequences, or reads, and an orthologous reference genome. When mapping T-rich reads to the reference genome, either a cytosine (C) or a thymine (T) in a read is considered a valid match for a cytosine in the reference genome. For A-rich reads, an adenine or a guanine is considered a valid match for a guanine in the reference genome. The mapping of reads to the reference genome by `rmapbs` is described below. If you choose to map reads with a different tool, make sure that your post-mapping files are appropriately formatted for the next components of the `methpipe` pipeline (necessary file formats for each step are covered in the corresponding sections). The default behavior of rmapbs is to assume that reads are T-rich and map accordingly. To change the mapping to suit A-rich reads, add the `-A` option.

Input and output file formats: We assume that the original data is a set of sequenced read files, typically as produced by Illumina sequencing. These are FASTQ format files, and can be quite large. After the reads are mapped, these files are not used by our pipeline. The reference genome should be a folder containing an individual FASTA (named like `*.fa`) file for each chromosome to maximize memory efficiency.

The mapped reads files (`*.mr` suffix) that result from the previous steps should consist of eight columns of data. The first six columns are the traditional components of a BED file (chromosome, start, end, read name, number of mismatches, strand), while the last two columns consist of sequence and quality scores respectively. These mapped reads files will be the input files for the following two `methpipe` components, `bsrate` and `methcounts`.

Decompressing and isolating paired-end reads: Sometimes paired-end reads are stored in the same FASTQ file. Because we treat these paired ends differently, they must be separated into two files and run through `rmapbs` with different parameters.

If your data is compressed as a Sequenced Read Archive, or SRA file, you can decompress and split paired-end reads into two files at the same time using `fastq-dump`, which is a program included in the `sra-toolkit` package, available for most unix systems. Below is an example of using `fastq-dump` to decompress and separate FASTQ data by end:

```
$ ./fastq-dump --split-3 Human_ESC.sra
```

If you have a FASTQ file not compressed in SRA format, you can split paired ends into two separate files by running the following commands:

```
$ sed -ne '1~8{N;N;N;p}' *.fastq > *_1.fastq
$ sed -ne '4~8{N;N;N;p}' *.fastq > *_2.fastq
```

Sequencing adaptors: These are a problem in any sequencing experiment with short fragments relative to the lengths of reads. `rmapbs` identifies sequences at the ends of reads greater than 10bp belonging to sequencing adaptors and converts them to Ns to avoid potential mapping problems.

Adaptor sequences must be supplied to `rmapbs` through the `-C` option. Keep in mind that if the adaptor sequence provided to you for the second paired end is displayed from 5' to 3', you will need to provide the reverse complement of the sequence to `rmapbs`.

Single-end reads: When working with data from a single-end sequencing experiment, you will have T-rich reads only. `rmapbs` expects T-rich reads as a default and so you do not have use the `-A` option to change mapping parameters. Execute the following command to map all of your single-end reads with `rmapbs`:

```
$ ./rmapbs -c hg18 -o Human_NHFF.mr Human_NHFF.fastq
```

Paired-end reads: When working with data from a paired-end sequencing experiment, you will have T-rich and A-rich reads. T-rich reads are often kept in files labeled with an "_1" and A-rich reads are often kept in files labeled with an "_2". T-rich reads are sometimes referred to as $5'$ reads or mate 1 and A-rich reads are sometimes referred to $3'$ reads or mate 2. We assume that the T-rich file and the A-rich contain the same number of reads, and each pair of mates occupy the same lines in their respective files. We will follow this convention throughout the manual and strongly suggest that you do the same. The program `rmapbs-pe` program is used to map T-rich reads and A-rich reads simultaneously. Run the following command to map two reads files from a paired-end sequencing experiment:

```
$ ./rmapbs-pe -c hg18 -o Human_ESC.mr Human_ESC_1.fastq Human_ESC_2.fastq
```

In brief, what happens internally in `rmapbs-pe` is as follows. `rmapbs-pe` finds candidate mapping locations for a T-rich mate with CG-wildcard mapping, and candidate mapping locations for the corresponding A-rich mate with AG-wildcard mapping. If two candidate mapping locations of the pair of mates are within certain distance in the same chromosome and strand and with correct orientation, the two mates are combined into a single read (after reverse complement of the A-rich mate), referred to as a fragment. The overlapping region between the two mates, if any, is included once, and the gap region between them, if any, is filled with Ns. The parameter `-L` to `rmapbs-pe` indicates the maximum size of fragments to allow to be merged. Here the fragment size is the sum of the read lengths at both ends, plus whatever distance is between them. So this is the length of the original molecule that was sequenced, excluding the sequencing adaptors. It is possible for a given read pair that the molecule was shorter than twice the length of the reads, in which case the ends of the mates will overlap, and so in the merged fragment will only be included once. Also, it is possible that the entire molecule was shorter than the length of even one of the mates, in which case the merged fragment will be shorter than either of the read ends. If the two mates cannot be merged because they are mapped to different chromosomes or different strand, or they are far away from each other, `rmapbs-pe` will throw each mate individually if its mapping position is unambiguous.

Mapping reads in a large file: Mapping reads often takes a while, and mapping reads from BS-seq takes even longer. It usually take quite a long time to map reads from a single large file with tens of millions of reads. If you have access to a cluster, one strategy is to launch multiple jobs, each working on a subset of reads simultaneously, and finally combine their output. I will typically map 3M reads at a time, and this takes at most 1.5GB of memory for the human genome and with 100nt reads. If all computing nodes can read the large input file (for example, through NFS), you may use the option `-T s` and `-N n` to instruct `rmapbs` or `rmapbs-pe` to map n reads starting from the s^{th} reads in the input file. For example, with following command

```
$ ./rmapbs -c hg18 -o Human_NHFF.mr Human_NHFF.fastq -T 1 -N 1000000
```

`rmapbs` will map the first million reads in the input file.

If each node can only access its local storage, dividing the set of reads to into k equal sized smaller reads files, and mapping these all simultaneously on multiple nodes, will make the mapping finish about k times faster. The unix `split` command is good for dividing the reads into smaller parts. The following BASH commands will take a directory named `reads` containing Illumina sequenced reads files, and split them into files containing at most 3M reads:

```
$ mkdir reads_split
$ for i in reads/*.txt; do \
    split -a 3 -d -l 12000000 ${i} reads_split/$(basename $i); done
```

Notice that the number of lines per split file is 12M, since we want 3M reads, and there are 4 lines per read. If you split the reads like this, you will need to "unsplit" them after the mapping is done. Not a problem, just use the `cat` command.

Alternative mappers: In addition to `rmapbs` described above, users may also wish to process raw reads using alternative mapping algorithms, including BSSeeker, which uses a three nucleotide alphabet strategy, and BSMAP, which allows for gaps during mapping. The program `to-mr` is used to convert the output from those mappers to the `*.mr` format used in our pipeline. To convert BSMAP mapped read file in .bam format, run

```
$ ./to-mr -o Human_NHFF.mr -m bsmap Human_NHFF.bam
```

where the option `-m` specifies that the original mapper is BSMAP. To obtain a list of alternative mappers supported by our converter, run `to-mr` without any options.

2.2 Merging libraries and removing duplicates

Before calculating methylation level, you should now remove read duplicates, or reads that were mapped to identical genomic locations. These reads are most likely the results of PCR over-amplication rather than true representations of distinct DNA molecules. The program `duplicate-remover` aims to remove such duplicates. It collects duplicate reads and/or fragments that have identical sequences and are mapped to the same genomic location (same chromosome, same start and end, and same strand), and chooses a random one to be the representative of the original DNA sequence.

`duplicate-remover` can take reads sorted by (chrom, start, end, strand). If the reads in the input file are not sorted, run the following sort command:

```
$ LC_ALL=C sort -k 1,1 -k 2,2n -k 3,3n -k 6,6 \
      -o Human_ESC.mr.sorted_start Human_ESC.mr
$ LC_ALL=C sort -k 1,1 -k 2,2n -k 3,3n -k 6,6 \
      -o Human_NHFF.mr.sorted_start Human_NHFF.mr
```

Next, execute the following command to remove duplicate reads:

```
$ ./duplicate-remover -S Human_ESC_dremove_stat.txt \
                       -o Human_ESC.mr.dremove Human_ESC.mr.sorted_start
$ ./duplicate-remover -S Human_NHFF_dremove_stat.txt \
                       -o Human_NHFF.mr.dremove Human_NHFF.mr.sorted_start
```

The duplicate-removal correction should be done on a per-library basis, i.e, one should pool all reads from multiple runs or lanes sequenced from the same library and remove duplicates. The reads from distinct libraries can be simply pooled without any correction as the reads from each library are originated from distinct DNA fragments. Please refer to 5 for recommended practices to organize a project with multiple runs and/or libraries.

4

2.3 Estimating bisulfite conversion rate

Unmethylated cytosines in DNA fragments are converted to uracils by sodium bisulfite treatment. As these fragments are amplified, the uracils are converted to thymines and so unmethylated Cs are ultimately read as Ts (barring error). Despite its high fidelity, bisulfite conversion of C to T does have some inherent failure rate, depending on the bisulfite kit used, reagent concentration, time of treatment, etc., and these factors may impact the success rate of the reaction. Therefore, the bisulfite conversion rate, defined as the rate at which unmethylated cytosines in the sample appear as Ts in the sequenced reads, should be measured and should be very high (*e.g.* > 0.99) for the experiment to be considered a success.

Measuring the bisulfite conversion rate this way requires some kind of control set of genomic cytosines not believed to be methylated. Three options are (1) to spike in some DNA known not to be methylated, such as a Lambda virus, (2) to use the human mitochondrial genome, which is known to be entirely unmethylated, or chloroplast genomes which are believed not to be methylated, or (3) to use non-CpG cytosines which are believed to be almost completely unmethylated in most mammalian cells. In general the procedure is to identify the positions in reads that correspond to these presumed unmethylated cytosines, then compute the ratio of C to (C + T) at these positions. If the bisulfite reaction was perfect, then this ratio should be very close to 1, and if there is no bisulfite treatment, then this ratio should be close to 0.

The program `bsrate` will estimate the bisulfite conversion rate in this way. Assuming method (3) from the above paragraph of measuring conversion rate at non-CpG cytosines in a mammalian methylome, the following command will estimate the conversion rate.

```
$ ./bsrate -c hg18 -o Human_ESC.bsrate Human_ESC.mr
$ ./bsrate -c hg18 -o Human_NHFF.bsrate Human_NHFF.mr
```

The `bsrate` program requires that the input be sorted so that reads mapping to the same chromosome are contiguous. The first several lines of the output might look like the following:

```
OVERALL CONVERSION RATE = 0.994141
POS CONVERSION RATE = 0.994166   832349
NEG CONVERSION RATE = 0.994116   825919
BASE PTOT  PCONV PRATE   NTOT   NCONV NRATE   BTHTOT BTHCONV BTHRATE ERR ALL    ERRRATE
1     8964  8813  0.9831  9024   8865  0.9823  17988  17678   0.9827  95  18083  0.0052
2     7394  7305  0.9879  7263   7183  0.9889  14657  14488   0.9884  100 14757  0.0067
3     8530  8442  0.9896  8323   8232  0.9890  16853  16674   0.9893  98  16951  0.0057
4     8884  8814  0.9921  8737   8664  0.9916  17621  17478   0.9918  76  17697  0.0042
5     8658  8596  0.9928  8872   8809  0.9929  17530  17405   0.9928  70  17600  0.0039
6     9280  9218  0.9933  9225   9177  0.9948  18505  18395   0.9940  59  18564  0.0031
7     9165  9117  0.9947  9043   8981  0.9931  18208  18098   0.9939  69  18277  0.0037
8     9323  9268  0.9941  9370   9314  0.9940  18693  18582   0.9940  55  18748  0.0029
9     9280  9228  0.9944  9192   9154  0.9958  18472  18382   0.9951  52  18524  0.0028
10    9193  9143  0.9945  9039   8979  0.9933  18232  18122   0.9939  66  18298  0.0036
```

The above example is based on a very small number of mapped reads in order to make the output fit the width of this page. The first thing to notice is that the conversion rate is computed separately for each strand. The information is presented separately because this is often a good way to see when some problem has occurred in the context of paired-end reads. If the conversion rate looks significantly different between the two strands, then we would go back and look for a mistake that has been made at an earlier stage in the pipeline. The first 3 lines in the output indicate the overall conversion rate, the conversion rate for positive strand mappers, and the conversion rate for negative strand mappers. The total number of nucleotides used (*e.g.* all C+T mapping over genomic non-CpG C's for method (3)) is given for positive and negative strand conversion rate computation, and if everything has worked up to this point these two numbers should be very similar. The 4th line gives column labels for a table showing conversion rate at each position in the reads. The labels PTOT, PCONV and PRATE give the total nucleotides used, the number converted, and the ratio of those two, for the positive-strand mappers. The corresponding numbers are also given for negative strand mappers (NTOT, NCONV, NRATE) and combined (BTH). The sequencing error rate is also shown for each

position, though this is an underestimate because we assume at these genomic sites any read with either a C or a T contains no error.

If you are using reads from an unmethylated spike-in or reads mapping to mitochondria, then there is an option to use all Cs, including those at CpG sites:

```
$ grep ^chrM Human_ESC.mr > Human_ESC.mr.chrM
$ ./bsrate -N -c chrM.fa -o Human_ESC.bsrate Human_ESC.mr.chrM
```

After completing bisulfite conversion rate analysis, remember to remove any control reads not naturally occurring in the sample (lambda virus, mitochondrial DNA from another organism, etc.) before continuing. The output from two different runs of `bsrate` can be merged using the program `merge-bsrate`.

2.4 Computing single-site methylation levels

The `methcounts` program takes the mapped reads and produces the methylation level at each genomic cytosine, with the option to produce only levels for CpG-context cytosines. While most DNA methylation exists in the CpG context, cytosines in other sequence contexts, such as CXG or CHH (where H denotes adenines, thymines, or cytosines and X denotes adenines or thymines) may also be methylated. Non-CpG methylation occurs most frequently in plant genomes and pluripotent mammalian cells such as embryonic stem cells. This type of methylation is asymmetric since the cytosines on the complementary strand do not necessarily have the same methylation status.

The input is in MappedRead format, and the reads should be sorted according to (chrom, end, start, strand). If your reads are not sorted, run:

```
$ LC_ALL=C sort -k 1,1 -k 3,3n -k 2,2n -k 6,6 \
        -o Human_ESC.mr.sorted_end_first Human_ESC.mr
```

Since `methcounts` can only take one input file, if you have multiple you can merge them using the `-m` option to the `sort` program:

```
$ LC_ALL=C sort -m -k 1,1 -k 3,3n -k 2,2n -k 6,6 \
        -o Human_ESC.mr.sorted_end_first Human_ESC.mr.1 Human_ESC.mr.2
```

Running `methcounts`: The methylation level for every cytosine site at single base resolution is estimated as a probability based on the ratio of methylated to total reads mapped to that loci. Because not all DNA methylation contexts are symmetric, methylation levels are produced for both strands and can be analyzed separately. To compute methylation levels at each cytosine site along the genome you can use the following command:

```
$ ./methcounts -c hg38 -o Human_ESC.meth \
               -S Human_ESC.methstats Human_ESC.mr
$ ./methcounts -c hg18 -o Human_NHFF.meth \
               -S Human_NHFF.methstats Human_NHFF.mr
```

The argument `-c` gives the filename of the genome sequence or the directory that contains one FASTA format file for each chromosome. By default `methcounts` identifies these chromosome files by the extension `.fa`. Importantly, the "name" line in each chromosome file must be the character > followed by the same name that identifies that chromosome in the mapped read output (the `.mr` files).

The output file contains one line per cytosine site. The first column is the chromosome. The second is the location of the cytosine. The 3rd column indicates the strand, which can be either + or −. The 4th column is the sequence context of that site, followed by an x if the site has mutated in the sample away from the reference genome. The 5th column is the estimated methylation level, equal to the number of Cs in reads at position corresponding to the site, divided by the sum of the Cs and Ts mapping to that position. The final column is number of reads overlapping with that site.

Note that because `methcounts` produces a file containing one line for every cytosine in the genome, the file can get quite large. For reference assembly mm10, the output is approximately 25GB. The `-n` option produces

methylation data for CpG context sites only, and for mm10 this produces an output file that is approximately 1GB. It is recommended that users allocate at least 8GB of memory when running `methcounts`.

To examine the methylation status of cytosines a particular sequence context, one may use the `grep` command to filter those lines based on the fourth column. For example, in order to pull out all cytosines within the CHG context, run the following:

```
$ grep CHG Human_ESC_All.meth > Human_ESC_CHG.meth
```

Our convention is to name `methcounts` output with all cytosines like `*_All.meth`, with CHG like `*_CHG.meth` and with CHH like `*_CHH.meth`.

Merging methcounts files from multiple replicates: When working with a BS-seq project with multiple replicates, you may first produce a methcounts output file for each replicate individually and assess the reproducibility of the methylation result by comparing different replicates. The `merge-methcounts` program is used to merge the those individual methcounts file to produce a single estimate that has higher coverage. Suppose you have the three methcounts files from three different biological replicates, `ASDF_R1/R1.meth`, `ASDF_R2/R2.meth` and `ASDF_R3/R3.meth`. To merge those individual methcounts files, execute

```
$ merge-methcounts ASDF_R1/R1.meth  ASDF_R2/R2.meth \
  ASDF_R3/R3.meth -o ASDF.meth
```

Computation of methylation level statistics The `levels` program computes statistics for the output of `methcounts`. Sample output is below. It computes the total fraction of cytosines covered, the fraction of cytosines that have mutated away from the reference, and coverage statistics for both CpGs and all cytosines. For CpG sites, coverage number reflects taking advantage of their symmetric nature and merging the coverage on both strands. For CpG coverage minus mutations, we remove the reads from CpG sites deemed to be mutated away from the reference. It also computes average methylation in three different ways, described in Schultz et al. (2012). This program should provide flexibility to compare methylation data with publications that calculate averages different ways and illustrate the variability of the statistic depending on how it is calculated.

```
SITES: 1000000
SITES COVERED: 566157
FRACTION MUTATED: 0.001257
FRACTION COVERED: 0.566157
MAX COVERAGE: 439
SYMMETRICAL CpG COVERAGE: 11.3228
SYMMETRICAL CpG COVERAGE (WHEN > 0): 15.3289
SYMMETRICAL CpG COVERAGE (minus mutations): 11.2842
SYMMETRICAL CpG COVERAGE (WHEN > 0) (minus mutations): 15.2768
MEAN COVERAGE: 3.31458
MEAN COVERAGE (WHEN > 0): 5.85452
METHYLATION LEVELS (CpG CONTEXT):
mean_meth 0.700166
w_mean_meth 0.667227
frac_meth 0.766211
METHYLATION LEVELS (CHH CONTEXT):
mean_meth 0.0275823
w_mean_meth 0.0184198
frac_meth 0.0146346
METHYLATION LEVELS (CXG CONTEXT):
mean_meth 0.0217537
w_mean_meth 0.0170535
```

```
frac_meth 0.00843068
METHYLATION LEVELS (CCG CONTEXT):
mean_meth 0.0211243
w_mean_meth 0.0187259
frac_meth 0.00630109
```

To run the `levels` program, execute

```
$ levels -o ASDF.levels ASDF.meth
```

3 Methylome analysis

The following tools will analyze much of the information about CpG's generated in previous steps and produce methylome wide profiles of various methylation characteristics. In the context of Methpipe, these characteristics consist of hypomethylated regions (HMRs), partially methylated regions (PMRs), differentially methylated regions between two methylomes (DMRs), regions with allele-specific methylation (AMRs), and hydroxymethylation.

3.1 Hypomethylated and hypermethylated regions (HMRs)

The distribution of methylation levels at individual sites in a methylome (either CpGs or non-CpG Cs) almost always has a bimodal distribution with one peak low (very close to 0) and another peak high (close to 1). In most mammalian cells, the majority of the genome has high methylation, and regions of low methylation are typically more interesting. These are called *hypo-methylated regions* (HMRs). In plants, most of the genome has low methylation, and it is the high parts that are interesting. These are called *hyper-methylated regions*. For stupid historical reasons in the Smith lab, we call both of these kinds of regions HMRs. One of the most important analysis tasks is identifying the HMRs, and we use the `hmr` program for this. The `hmr` program uses a hidden Markov model (HMM) approach using a Beta-Binomial distribution to describe methylation levels at individual sites while accounting for the number of reads informing those levels. `hmr` automatically learns the average methylation levels inside and outside the HMRs, and also the average size of those HMRs.

Requirements on the data: We typically like to have about 10x coverage to feel very confident in the HMRs called in mammalian genomes, but the method will work with lower coverage. The difference is that the boundaries of HMRs will be less accurate at lower coverage, but overall most of the HMRs will probably be in the right places if you have coverage of 5-8x (depending on the methylome). Boundaries of these regions are totally ignored by analysis methods based on smoothing or using fixed-width windows.

Typical mammalian methylomes: Running `hmr` requires a file of methylation levels formatted like the output of the `methcounts` program (as described above). For calling HMRs in mammalian methylomes, we suggest only considering the methylation level at CpG sites, as the level of non-CpG methylation is not usually more than a few percent. You can do this by filtering the `methcounts` output with `grep` (as below) and running `hmr` with the CpG-only file:

```
$ grep CpG Human_ESC_All.meth > Human_ESC.meth
$ ./hmr -o Human_ESC.hmr Human_ESC.meth
```

The output will be in BED format, and the indicated strand (always positive) is not informative. The name column in the output will just assign a unique name to each HMR. Each time the `hmr` is run it requires parameters for the HMM to use in identifying the HMRs. We usually train these HMM parameters on the data being analyzed, since the parameters depend on the average methylation level and variance of methylation level; the variance observed can also depend on the coverage. However, in some cases it might be desirable to use the parameters trained on one data set to find HMRs in another. The option `-p` indicates a file in which the trained parameters are written, and the argument `-P` indicates a file containing parameters (as produced with the `-p` option on a previous run) to use:

```
$ ./hmr -p Human_ESC.hmr.params -o Human_ESC.hmr Human_ESC.meth
$ ./hmr -P Human_ESC.hmr.params -o Human_NHFF_ESC_params.hmr Human_NHFF.meth
```

In the above example, the parameters were trained on the ESC methylome, stored in the file `Human_ESC.hmr.params` and then used to find HMRs in the NHFF methylome. This is useful if a particular methylome seems to have very strange methylation levels through much of the genome, and the HMRs would be more comparable with those from some other methylome if the model were not trained on that strange methylome.

Plant (and similar) methylomes: The plant genomes, exemplified by *A. thaliana*, are devoid of DNA methylation by default, with genic regions and transposons being hyper-methylated, which we termed HyperMRs to stress their difference from *hypo-methylated regions* in mammalian methylomes. DNA methylation in plants has been associated with expression regulation and transposon repression, and therefore characterizing HyperMRs is of much biological relevance. In addition to plants, hydroxymethylation tends to appear in a small fraction of the mammalian genome, and therefore it makes sense to identify hyper-hydroxymethylated regions.

The first kind of HyperMR analysis involves finding continuous blocks of hyper-methylated CpGs with the `hmr` program. Since `hmr` is designed to find hypo-methylated regions, one needs first to invert the methylation levels in the `methcounts` output file as follows:

```
$ awk '{$5=1-$5; print $0}' Col0.meth > Col0_inverted.meth
```

Next one may use the `hmr` program to find "valleys" in the inverted Arabidopsis methylome, which are the hyper-methylated regions in the original methylome. The command is invoked as below

```
$ ./hmr -o Col0.hmr Col0_inverted.meth
```

This kind of HyperMR analysis produces continuous blocks of hyper-methylated CpGs. However in some regions, intragenic regions in particular, such continuous blocks of hyper-methylated CpGs are separated by a few unmethylated CpGs, which have distinct sequence preference when compared to those CpGs in the majority of unmethylated genome. The blocks of hyper-methylated CpGs and gap CpGs together form composite HyperMRs. The `hypermr` program, which implements a three-state HMM, is used to identify such HyperMRs. Suppose the `methcounts` output file is `Col0_Meth.bed`, to find HyperMRs from this dataset, run

```
$ ./hypermr -o Col0.hypermr Col0.meth
```

The output file is a 6-column BED file. The first three columns give the chromosome, starting position and ending position of that HyperMR. The fourth column starts with the "hyper:", followed by the number of CpGs within this HyperMR. The fifth column is the accumulative methylation level of all CpGs. The last column indicates the strand, which is always +.

Lastly, it is worth noting that plants exhibit significantly more methylation in the non-CpG context, and therefore inclusion of non-CpG methylation in the calling of hyper-methylated regions could possibly be informative. We suggest separating each cytosine context from the `methcounts` output file as illustrated in the previous section (via `grep`) and calling HyperMRs seperately for each context.

Partially methylated regions (PMRs): The `hmr` program also has the option of directly identifying partially methylated regions (PMRs), not to be confused with partially methylated domains (see below). These are contiguous intervals where the methylation level at individual sites is close to 0.5. This should also not be confused with regions that have allele-specific methylation (ASM) or regions with alternating high and low methylation levels at nearby sites. Regions with ASM are almost always among the PMRs, but most PMRs are not regions of ASM. The `hmr` program is run with the same input but a different optional argument to find PMRs:

```
$ ./hmr -partial -o Human_ESC.pmr Human_ESC.meth
```

Giant HMRs observed in cancer samples (AKA PMDs): Huge genomic blocks with abnormal hypomethylation have been extensively observed in human cancer methylomes and more recently in extraembryonic tissues like the placenta. These domains are characterized by enrichment in intergenic regions or Lamina associated domains (LAD), which are usually hypermethylated in normal tissues. Partially methylated domains (PMDs) are not homogeneously hypomethylated as in the case of HMRs, and contain focal hypermethylation at specific sites. Hidden Markov Models can also identify these larger domains. The program `pmd` is provided for their identification, and can be run as follows:

```
$ ./pmd -o Human_ESC.pmd Human_ESC.meth
```

3.2 Differential methylation between two methylomes

If you are working with more than one methylome, it may be of interest to you to identify regions between your methylomes that have significantly different levels of methylation. To do this, use the programs `methdiff` and `dmr`. Run `methdiff` first since its output serves as the input for `dmr`. Since methylation differences are assessed on a per CpG basis, the methylomes being compared must come from the same genomes. Otherwise, comparisons will not be between orthologous CpG's. If you would like to compare methylomes from different genomes (i.e. human and chimp methylomes), you must first convert the CpG coordinates for one species into their orthologous coordinates for the other species. Additionally, `methdiff` and `dmr` can only compare two methylomes at a time. Each of these programs is explained in more detail in the subsections below.

3.2.1 Differential methylation scores

The program `methdiff` produces a differential methylation score for each CpG in a methylome. This score indicates the probability that the CpG is significantly less methylated in one methylome than the other. The inputs for `methdiff` are the output of `methcounts` for each of the two methylomes being analyzed. The following command calculates differential methylation scores across two methylomes using the `methdiff` component of Methpipe.

```
$ ./methdiff -o Human_ESC_NHFF.methdiff \
             Human_ESC.meth Human_NHFF.meth
```

So in the output file `Human_ESC_NHFF.methdiff` the 5th column indicates the probability that the methylation level at each given site is lower in `Human_NHFF.meth` than in `Human_ESC.meth`. For the other direction, you can either swap the order of those two input files, or just subtract the probability from 1.0. The method used is due to Altham (1971) [**?**], and is like a one-directional version of Fisher's exact test. The fourth column of the output file contains the cytosine's context, and then the number of reads methylated and total coverage for each of the two samples being compared at that site.

3.2.2 Differentially methylated regions (DMRs)

Once differential methylation scores have been calculated, the program `dmr` can be used to identify differentially methylated regions, or DMRs. DMRs are regions where differential methylation scores indicate there are many CpGs with a high probability of being differentially methylated between the two methylomes. `dmr` uses HMR data from the two methylomes and identifies a DMR wherever an HMR exists in one methylome but not the other. It writes the DMRs into two files: one with the HMR in one methylome and another with the HMR in the other. It also writes the total number of CpGs in the DMR and the number of significantly different CpG sites. The following command finds DMRs using the `dmr` component of `methpipe`:

```
$ ./dmr Human_ESC_NHFF.methdiff Human_ESC.hmr \
     Human_NHFF.hmr DMR_ESC_lt_NHFF DMR_NHFF_lt_ESC
```

3.3 Allele-specific methylation

Allele-specific methylation (ASM) occurs when the same cytosine is differentially methylated on the two alleles of a diploid organism. ASM is a major mechanism of genomic imprinting, and aberrations can lead to disease. Included in `methpipe` are three tools to analyze ASM: `allelicmeth`, `amrfinder`, and `amrtester`. All of these programs calculate the probability of ASM in a site or region by counting methylation on reads and analyzing the dependency between adjacent CpGs, and therefore it is recommended that any samples analyzed have at least $10\times$ coverage and 100bp reads for the human genome.

3.3.1 Single-site ASM scoring

The program `allelicmeth` calculates allelic-specific methylation scores for each CpG site. Input files should be the mapped reads files (`.mr` suffix) produced previously in the mapping step. In the output file, each row represents a CpG pair made by any CpG and its previous CpG, the first three columns indicate the positions of the CpG site, the fourth column is the name including the number of reads covering the CpG pair, the fifth column is the score for ASM, and the last four columns record the number of reads of four different methylation combinations of the CpG pair: methylated methylated (mm), methylated unmethylated (mu), unmethylated methylated (um), or unmethylated unmethylated (uu). The following command will calculate allele-specific methylation scores using the `allelicmeth` component of Methpipe.

```
$ ./allelicmeth -c hg18 -o Human_ESC.allelicmeth Human_ESC.mr
```

3.3.2 Allelically methylated regions (AMRs)

The method described here was introduced in [**?**]. The program `amrfinder` scans the genome using a sliding window to identify AMRs. For a genomic interval, two statistical models are fitted to the reads mapped, respectively. One model (single-allele model) assumes the two alleles have the same methylation state, and the other (two-allele model) represents different methylation states for the two alleles. Comparing the likelihood of the two models, the interrogated genomic interval may be classified as an AMR. The input files must be in the "epiread" format, which consists of three columns. The first column is the chromosome of the read, the second column is the numbering order of the first CpG in the read, and the last column is the CpG only sequence of the read.

You can use the program `methstates` to convert the mapped reads file to the epiread format as below:

```
$ ./methstates -c hg18 -o Human_ESC.epiread Human_ESC.mr
```

The following command shows an example to run the program `amrfinder`.

```
$ ./amrfinder -o Human_ESC.amr -c hg18 Human_ESC.epiread
```

There are several options for running `amrfinder`. The `-b` switches from using a likelihood ratio test to BIC as the criterion for calling an AMR. The `-i` option changes the number of iterations used in the EM procedure when fitting the models. The `-w` option changes the size of the sliding window, which is in terms of CpGs. The default of 10 CpGs per window has worked well for us. The `-m` indicates the minimum coverage per CpG site required for a window to be tested as an AMR. The default requires 4 reads on average, and any lower will probably lead to unreliable results. AMRs are often fragmented, as coverage fluctuates, and spacing between CpGs means their linkage cannot be captured by the model. The `-g` parameter is used to indicate the maximum distance between any two identified AMRS; if two are any closer than this value, they are merged. The default is 1000, and it seems to work well in practice, not joining things that appear as though they should be distinct. In the current version of the program, at the end of the procedure, any AMRs whose size in terms of base-pairs is less than half the "gap" size are eliminated. This is a hack that has produced excellent results, but will eventually be eliminated (hopefully soon).

Finally, the `-C` parameter specifies the critical value for keeping windows as AMRs, and is only useful when the likelihood ratio test is the used; for BIC windows are retained if the BIC for the two-allele model is less than that for the single-allele model. `amrfinder` calculates a false discovery rate to correct for multiple testing, and therefore most p-values that pass the test will be significantly below the critical value. The `-h` option produces FDR-adjusted

p-values according to a step-up procedure and then compares them directly to the given critical value, which allows further use of the p-values without multiple testing correction. The `-f` omits multiple testing correction entirely by not applying a correction to the p-values or using a false discovery rate cutoff to select AMRs.

In addition to `amrfinder`, which uses a sliding window, there is also the `amrtester` program, which tests for allele-specific methylation in a given set of genomic intervals. The program can be run like this:

```
$ ./amrtester -o Human_ESC.amr -c hg18 intervals.bed Human_ESC.epiread
```

This program works very similarly to `amrfinder`, but does not have options related to the sliding window. This program outputs a score for each input interval, and when the likelihood ratio test is used, the score is the p-value, which can easily be filtered later.

3.4   Consistent estimation of hydroxymethylation and methylation levels

If you are interested in estimating hydroxymethylation level and have any two of Tet-Assisted Bisulfite sequencing (TAB-seq), oxidative bisulfite sequencing (oxBS-seq) and BS-seq data available, you can use `mlml` [**?**] to perform consistent and simultaneous estimation.

The input file format could be the default `methcounts` output format described in Section 2.4, or BED format file with 6 columns as the example below:

```
chr1     3001345 3001346 CpG:9   0.777777777778  +
```

Here the fourth column indicates that this site is a CpG site, and the number of reads covering this site is 9. The fifth column is the methylation level of the CpG site, ranging from 0 to 1. Note that all input files must be sorted. Assume you have three input files ready: `meth_BS-seq.meth`, `meth_oxBS-seq.meth` and `meth_Tab-seq.meth`. The following command will take all the inputs:

```
$ ./mlml -v -u meth_BS-seq.meth -m meth_oxBS-seq.meth \
        -h meth_Tab-seq.meth -o result.txt
```

If only two types of input are available, e.g. `meth_BS-seq.meth` and `meth_oxBS-seq.meth`, then use the following command:

```
$ ./mlml -u meth_BS-seq.meth -m meth_oxBS-seq.meth \
        -o result.txt
```

In some cases, you might want to specify the convergence tolerance for EM algorithm. This can be done through `-t` option. For example:

```
$ ./mlml -u meth_BS-seq.meth -m meth_oxBS-seq.meth \
        -o result.txt -t 1e-2
```

This command will make the iteration process stop when the difference of estimation between two iterations is less than 10^{-2}. The value format can be scientific notation, e.g. 1e-5, or float number, e.g. 0.00001.

The output of `mlml` is tab-delimited format. Here is an example:

```
chr11 15 16 0.166667 0.19697 0.636364 0
chr12 11 12 0.222222 0 0.777778 2
```

The columns are chromosome name, start posistion, end position, 5-mC level, 5-hmC level, unmethylated level and number of conflicts. To calculate the last column, a binomial test is performed for each input methylation level (can be 2 or 3 in total depending on parameters). If the estimated methylation level falls out of the confidence interval calculated from input coverage and methylation level, then such event is counted as one conflict. It is recommended to filter estimation results based on the number of conflicts; if more conflicts happens on one site then it is possible that information from such site is not reliable.

3.5 Computing average methylation level in a genomic interval

One of the most common analysis tasks is to compute the average methylation level through a genomic region. The `roimethstat` program accomplishes this. It takes a sorted `methcounts` output file and a sorted BED format file of genomic "regions of interest" (hence the "roi" in `roimethstat`). If either file is not sorted by (chrom,end,start,strand) it can be sorted using the following command:

```
$ LC_ALL=C sort -k 1,1 -k 3,3n -k 2,2n -k 6,6 \
         -o regions_ESC.meth.sorted regions_ESC.meth
```

From there, `roimethstat` can be run as follows:

```
$ ./roimethstat -o regions_ESC.meth regions.bed Human_ESC.meth.sorted
```

The output format is also 6-column BED, and the score column now takes the average methylation level through the interval, weighted according to the number of reads informing about each CpG or C in the methylation file. The 4th, or "name" column encodes several other pieces of information that can be used to filter the regions. The original name of the region in the input regions file is retained, but separated by a colon (:) are, in the following order, (1) the number of CpGs in the region, (2) the number of CpGs covered at least once, (3) the number of observations in reads indicating in the region that indicate methylation, and (4) the total number of observations from reads in the region. The methylation level is then (3) divided by (4). Example output might look like:

```
chr1  3011124  3015902  REGION_A:18:18:105:166  0.63253  +
chr1  3015904  3016852  REGION_B:5:5:14:31  0.451613  +
chr1  3017204  3017572  REGION_C:2:2:2:9  0.222222  -
chr1  3021791  3025633  REGION_D:10:10:48:73  0.657534  -
chr1  3026050  3027589  REGION_E:2:4:4:32:37  0.864865  -
```

Clearly if there are no reads mapping in a region, then the methylation level will be undefined. By default `roimethstat` does not output such regions, but sometimes they are helpful, and using the `-P` flag will force `roimethstat` to print these lines in the output (in which case every line in the input regions will have a corresponding line in the output).

3.6 Computing methylation entropy

The concept of Entropy was introduced into epigenetics study to characterize the randomness of methylation patterns over several consecutive CpG sites [**?**]. The `methentropy` program processes epireads and calculates the methylation entropy value in sliding windows of specified number of CpGs. Two input files are required, including the directory containing the chromosome fasta files, and an epiread file as produced by `methstates` program. Use the `-w` option to specify the desired number of CpGs in the sliding window; if unspecified, the default value is 4. In cases where symmetric patterns are considered the same, specify option `-F`, this will cause the majority state in each epiread to be forced into "methylated", and the minority to "unmethylated". The processed epireads will then be used for entropy calculation. To run the program, type command:

```
$ ./methentropy -w 5 -v -o Human_ESC.entropy hg18 Human_ESC.epiread
```

The output format is the same as `methcount` output. The first 3 columns indicate the genomic location of the center CpG in each sliding window, the 5th column contains the entropy values, and the 6th column shows the number of reads used for each sliding window. Below is an output example.

```
chr1 483 + CpG 2.33914 27
chr1 488 + CpG 2.05298 23
chr1 492 + CpG 1.4622 24
chr1 496 + CpG 1.8784 35
```

3.7 Notes on data quality

The performance of our tools to identify higher-level methylation features (HMR, HyperMR, PMDs and AMR) depends on the underlying data quality. One major factor is coverage. Based on our experience, HMR detection using our method is acceptable above 5x coverage, and we recommend 10x for reliable results. Our method for identifying HyperMRs is similar to HMR-finding method, and the above statement holds. The required coverage therefore is even lower since the PMD-finding method internally works by accumulating CpGs in fixed-length bins. We feel ~3x is sufficient. The AMR method depends on both coverage and read length. Datasets with read length around 100bp and mean coverage above 10x are recommended for the AMR method. Another important measure of data quality is the bisulfite conversion rate. Since most datasets have pretty good bisulfite conversion rate (above 0.95), our tools does not explicitly correct fo the conversion rate.

4 Methylome visualization

4.1 Creating UCSC Genome Browser tracks

To view the methylation level or read coverage at individual CpG sites in a genome browser, one needs to create a `bigWig` format file from a `*_.meth` file, which is the output of the `methcounts` program. A methcounts file would look like this:

```
chr1 468 469 CpG:30 0.7 +
chr1 470 471 CpG:29 0.931034 +
chr1 483 484 CpG:36 0.916667 +
chr1 488 489 CpG:36 1 +
```

The first 3 columns shows the physical location of each CpG sites in the reference genome. The number in the 4th column indicates the coverage at each CpG site. The methylation level at individual CpG sites can be found in the 5th column. To create methylation level tracks or read coverage tracks, one can follow these steps:

1. Download the `wigToBigWig` program from UCSC genome browser's directory of binary utilities (`http://hgdownload.cse.ucsc.edu/admin/exe/`).

2. Use the `fetchChromSizes` script from the same directory to create the `*.chrom.sizes` file for the UCSC database you are working with (e.g. hg19). Note that this is the file that is referred to as `hg19.chrom.sizes` in step 3.

3. To create a `bw` track for methylation level at single CpG sites, convert the methcounts file to bed format using:

   ```
   $ awk '{print $1 "\t" $2 "\t" $2+1 "\t" $4":"$6 "\t" $5 "\t" $3}' \
         Human_ESC.meth > Human_ESC.meth.bed
   ```

4. To create a `bw` track from the bed format methcounts output, modify and use the following command:

   ```
   $ cut -f 1-3,5 Human_ESC.meth.bed | \
        wigToBigWig /dev/stdin hg19.chrom.sizes Human_ESC.meth.bw
   ```

   To create a `bw` track for coverage at single CpG sites, modify and use the following command:

   ```
   $ tr ':' '[Ctrl+v Tab]' < Human_ESC.meth.bed | cut -f 1-3,5 | \
        wigToBigWig /dev/stdin hg19.chrom.sizes Human_ESC.reads.bw
   ```

Note that if the `wigToBigWig` or `fetchChromSizes` programs are not executable when downloaded, do the following:

```
$ chmod +x wigToBigWig
$ chmod +x fetchChromSizes
```

You might also want to create `bigBed` browser tracks for HMRs, AMRs, PMDs, or DMRs. To do so, follow these steps:

1. Download the `bedToBigBed` program from the UCSC Genome Browser directory of binary utilities (`http://hgdownload.cse.ucsc.edu/admin/exe/`).

2. Use the `fetchChromSizes` script from the same directory to create the `*.chrom.sizes` file for the UCSC database you are working with (e.g. hg19). Note that this is the file that is referred to as `hg19.chrom.sizes` in step 3.

3. Modify and use the following commands: For `*_HMR.bed` files with non-integer score in their 5th column, one needs to round the score to integer value, for example:

```
$ awk -v OFS="\t" '{print $1,$2,$3,$4,int($5)}' Human_ESC.hmr \
        > Human_ESC.hmr.rounded

 bedToBigBed Human_ESC.hmr.rounded hg19.chrom.sizes Human_ESC.hmr.bb
```

In the above command, since the HMRs are not stranded, we do not print the 6th column. Keeping the 6th column would make all the HMRs appear as though they have a direction – but it would all be the + strand.

```
$ bedToBigBed Human_ESC.hmr hg19.chrom.sizes Human_ESC.hmr.bb
```

5 Organizing projects with multiple replicates and/or libraries

Often to build a methylome for a particular population of cells one obtains multiple biological replicates. For each of these biological replicates, one may construct more than one bisulfite sequencing library, and these correspond to technical replicates. Finally, a given library might be sequenced multiple times, and a given sequencing run might produce one or more files of reads (for example, corresponding to different tiles or lanes on an Illumina sequencer). We organize data to reflect these aspects of the experiment.

Biological replicates Distinct biological replicates refer to distinct cells from which DNA was extracted. Typically these are obtained form different individuals or different cell cultures. If one is building the "refernce" methylome for a given cell type, then eventually these different biological replicates might be merged. However, the purpose of doing the replicates separately is so that the biological variation can be understood. Therefore, certain tests would usually be done prior to combining the different biological replicates. Fortunately, the methylomes corresponding to different biological replicates can be done at much lower coverage, because testing for biological variation does not require conducting all the kinds of analyses we would want to conduct on the reference methylome. Distinct biological replicates are necessarily distinct technical replicates, and so to combine them we would follow the same procedure as combining technical replicates (see below). If the name of our methylome is `ASDF` then we would organize our biological replicates in a directory structure like this:

```
./ASDF/ASDF_R1/
./ASDF/ASDF_R2/
./ASDF/ASDF_R3/
```

The properties that we would typically associate with a biological replicate include the average methylation levels through different parts of the genome. When comparing biological replicates we might observe some differences in average methylation levels through promoters, for example. We would hope and expect these to be minimal, resulting either from noise due to sampling (when coverage is low) or variation associated with genotype.

Distinct libraries   The distinct libraries constitute technical replicates. Each technical replicate may have a different distribution of fragment lengths (as measured using paired-end information or amount of adaptors present in reads), and possibly a different bisulfite conversion rate. Despite these differences, we should generally see similar methylation levels. If we compare methylation levels between methylation profiles from different library preparations, using a defined set of intervals, we would expect to see very little difference. The difference we do observe should reflect noise due to sampling in regions with low coverage. Because methylation data from two different libraries reflects different actual molecules, combining data from different libraries is easy. The methylation levels in `methcounts` files can be combined by simply merging the individual `methcounts` files. Similarly, the `methstates` files can be concatenated and the re-sorted. Within a biological replicate, we organize libraries as follows.

```
./ASDF/ASDF_R1/ASDF_R1_L1/
./ASDF/ASDF_R1/ASDF_R1_L2/
./ASDF/ASDF_R1/ASDF_R1_L3/
```

Reads files within a library   For a given library, having a simple pipeline is facilitated by having a uniform scheme for naming the data files. The following example illustrates files corresponding to 4 sets of reads, two from paired-end sequencing and two from single-end sequencing:

```
./ASDF/ASDF_R1/ASDF_R1_L1/1_1.fq
./ASDF/ASDF_R1/ASDF_R1_L1/1_2.fq
./ASDF/ASDF_R1/ASDF_R1_L1/2_1.fq
./ASDF/ASDF_R1/ASDF_R1_L1/2_2.fq
./ASDF/ASDF_R1/ASDF_R1_L1/3.fq
./ASDF/ASDF_R1/ASDF_R1_L1/4.fq
```

The files `1_1.fq` and `1_2.fq` are corresponding left and right mates files. These must be processed together so that the corresponding mates can be joined. The file `3.fq` is from a single-end sequencing run. We recommend using symbolic links to set up these filenames, and to keep subdirectories corresponding to the individual sequencing runs for the same library.

```
./ASDF/ASDF_R1/ASDF_R1_L1/Run1/s_1_1_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/Run1/s_1_2_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/Run2/LID12345_NoIndex_L001_R1_001.fastq
./ASDF/ASDF_R1/ASDF_R1_L1/Run2/LID12345_NoIndex_L001_R2_001.fastq
./ASDF/ASDF_R1/ASDF_R1_L1/Run3/s_1_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/Run3/s_2_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/1_1.fq --> Run1/s_1_1_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/1_2.fq --> Run1/s_1_2_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/2_1.fq --> Run2/LID12345_NoIndex_L001_R1_001.fastq
./ASDF/ASDF_R1/ASDF_R1_L1/2_2.fq --> Run2/LID12345_NoIndex_L001_R2_001.fastq
./ASDF/ASDF_R1/ASDF_R1_L1/3.fq --> Run3/s_1_sequence.txt
./ASDF/ASDF_R1/ASDF_R1_L1/4.fq --> Run3/s_2_sequence.txt
```

This kind of organization accomplishes four things: (1) it keeps read files from different runs separate, (2) it provides a place to keep metadata for the different runs, (3) it allows the reads files to exist with their original names which might not be different between runs, and (4) it ensures a simple naming scheme for the data files needed by our pipeline.

The results directories   For the example scheme we have been describing, we would organize the results in the following directories.

```
./ASDF/ASDF_R1/ASDF_R1_L1/results/
./ASDF/ASDF_R1/ASDF_R1_L2/results/
./ASDF/ASDF_R1/results/
```

```
./ASDF/ASDF_R2/ASDF_R2_L1/results/
./ASDF/ASDF_R2/ASDF_R2_L2/results/
./ASDF/ASDF_R2/results/
./ASDF/ASDF_R3/ASDF_R3_L1/results/
./ASDF/ASDF_R3/ASDF_R3_L2/results/
./ASDF/ASDF_R3/results/
./ASDF/results/
```

6 Auxiliary tools

6.1 Count number of lines in a big file

When working with next-generation sequencing data, researchers often handle very large files, such as FASTQ files containing raw reads and `*.mr` files containing mapped reads. `lc_approx` is an auxiliary tool designed to approximate the number of lines in a very large file by counting the number of lines in a small, randomly chosen chunk from the big file and scaling the estimate by file size. For example, in order to estimate the number of reads in a FASTQ file `s_1_1_sequence.fq`, run

```
$ lc_approx  s_1_1_sequence.fq
```

It will return the approximate number of lines in this file and by dividing the above number by 4, you get the approximate number of reads in that file. The `lc_approx` can be hundreds of times faster than the unix tool `wc -l`.

6.2 Automating methylome analysis

Two bash scripts have been provided to perform quick and consistent analysis on the individual library level and aggregated biological replicate level. For each library in a project, `library` sorts, removes duplicate reads, calculates the bisulfite conversion rate, runs `methcounts`, and runs `levels`. There is also an option to perform `methcounts` analysis on all cytosines, which can be toggled by uncommenting the command in the script. For RRBS data, the `duplicate-remover` command should be commented.

To run `library`, navigate to the library's results directory – the structure should look something like

```
./ASDF/ASDF_R1/ASDF_R1_L1/results_mm9/
```

and should contain a single mapped reads file. To run, specify the directory where `methpipe` binaries can be found, the directory where chromosome files can be found, and the base name of the library or replicate you are working on:

```
library /home/user/Desktop/methpipe/trunk/bin \
        /home/user/Desktop/mm9_chroms ASDF_R1_L1
```

When `library` has been run on all individual libraries, the results can be merged to produce results for the biological replicate using `merge-methylomes`. This merges bisulfite conversion rate and methcounts statistics, and then uses them to generate HMRs, PMRs, and AMRs. To run `merge-methylomes`, navigate to the biological replicate's results directory – the structure should look something like

```
./ASDF/ASDF_R1/results_mm9/
```

and should be empty when the program starts. To run, specify the same command line arguments as above, with the biological replicate as the base name:

```
biorep /home/user/Desktop/methpipe/trunk/bin \
        /home/user/Desktop/mm9_chroms ASDF_R1 L
```

In some cases, it is useful to merge biological replicates rather than libraries. In this case, the fourth parameter of `merge-methylomes` should be R.

All run configurations for the programs in these scripts are consistent with data in methbase, and therefore direct comparison is appropriate. These tools should provide a convenient, consistent workflow for researchers to quickly analyze and compare their methylomes with those made publically available in methbase.

6.3 Mapping methylomes between species

Mapping methylomes between species builds on the the `liftOver` tool provided by UCSC Genome Browser `http://genome.ucsc.edu/cgi-bin/hgLiftOver`. However it is time consuming to directly convert large `methcounts` output files containing millions of CpG sites. We therefore provide the `fastLiftOver` tool.

 `fastLiftOver` requires its own chain files which are generated with the UCSC `liftOver` program. To generate `fastLiftOver` chain file from the mouse mm9 genome to the human hg19 genome, you need a BED file (CpGs-mm9.bed) with all CpG sites in the mouse genome and the `liftOver` chain file (mm9ToHg19.over.chain) from the UCSC Genome Browser website. The BED file of mm9 CpG sites follows the format below, with the fourth column composed of the chromosome name and the starting location.

```
chr1 3000573 3000574 chr1:3000573
chr1 3000725 3000726 chr1:3000725
chr1 3000900 3000901 chr1:3000900
...
```

First, we will use the UCSC `liftOver` program to map the mouse CpG sites to the human reference genome:

```
liftOver CpGs-mm9.bed mm9ToHg19.over.chain CpGs-mm9Tohg19.bed CpGs-mm9Tohg19-left.bed
```

Next, we need to sort the file `CpGs-mm9Tohg19.bed`, that contains CpG sites mapped to the human reference genome

```
export LC_ALL=C;
sort -k1,1 -k2,2n  CpGs-mm9Tohg19.bed -o tmpfile
mv tmpfile CpGs-mm9Tohg19.bed
```

The file CpGs-mm9Tohg19.bed will be the `fastLiftOver` chain file for mapping mouse methcounts output files to the human reference genome. For example, if we want to map the mouse methcounts file (Mouse_ESC.meth) to the human reference genome, we may run the following command:

```
fastLiftOver -f Mouse_ESC.meth -i CpGs-mm9Tohg19.bed -t Mouse_ESC-hg19.meth -v
```

 It is possible that multiple CpG sites in the mouse reference genome are mapped to the same location in the human reference genome, therefore the many-to-one conversion problem. In this situation, we may use the average methylation level weighted by coverage. Please checkout `https://github.com/songqiang/MethPipe/blob/master/run-fastLiftOver.sh` for example scripts that go over the whole procedure.